

sentiment-analysis-for-tweets

December 5, 2023

0.1 Hate Text detection

Dataset: <https://www.kaggle.com/datasets/arkhoshghalb/twitter-sentiment-analysis-hatred-speech>

1 Importing Required Libraries

```
[51]: import pandas as pd
import numpy as np
import re
import seaborn as sns
import matplotlib.pyplot as plt
from matplotlib import style
style.use('ggplot')
from nltk.tokenize import word_tokenize
from nltk.stem import WordNetLemmatizer
from nltk.corpus import stopwords
import nltk
nltk.download('stopwords')
nltk.download('punkt')
nltk.download('wordnet')
stop_words = set(stopwords.words('english'))
from wordcloud import WordCloud
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, classification_report, \
    confusion_matrix, ConfusionMatrixDisplay
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data] Package wordnet is already up-to-date!
```

```
[52]: tweet_df = pd.read_csv('/content/train.csv')
      tweet_df.head()
```

```
[52]:   id  label          tweet
0    1     0  @user when a father is dysfunctional and is s...
1    2     0  @user @user thanks for #lyft credit i can't us...
2    3     0                bihday your majesty
3    4     0  #model    i love u take with u all the time in ...
4    5     0                factsguide: society now      #motivation
```

```
[53]: tweet_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 31962 entries, 0 to 31961
Data columns (total 3 columns):
 #   Column  Non-Null Count  Dtype
---  -
0    id      31962 non-null    int64
1    label    31962 non-null    int64
2    tweet    31962 non-null    object
dtypes: int64(2), object(1)
memory usage: 749.2+ KB
```

```
[54]: print(tweet_df['tweet'].iloc[0], "\n")
```

```
@user when a father is dysfunctional and is so selfish he drags his kids into
his dysfunction.      #run
```

2 Data Preprocessing

- Removing links, mails, more than one spaces
- Tokenization using NLTK
- Removing stop words

```
[55]: def data_processing(tweet):
      tweet = tweet.lower()
      tweet = re.sub(r"https\S+|www\S+http\S+", '', tweet, flags = re.MULTILINE)
      tweet = re.sub(r'\@w+|\#', '', tweet)
      tweet = re.sub(r'[\w\s]', '', tweet)
      tweet = re.sub(r'ð', '', tweet)
      tweet_tokens = word_tokenize(tweet)
      filtered_tweets = [w for w in tweet_tokens if not w in stop_words]
      return " ".join(filtered_tweets)
```

```
[56]: tweet_df.tweet = tweet_df['tweet'].apply(data_processing)
      tweet_df = tweet_df.drop_duplicates('tweet')
```

- Lemmatization Using NLTK

```
[57]: lemmatizer = WordNetLemmatizer()
def lemmatizing(data):
    tweet = [lemmatizer.lemmatize(word) for word in data]
    return tweet
```

```
[58]: tweet_df['tweet'] = tweet_df['tweet'].apply(lambda x: lemmatizing(x))
```

```
[59]: print(tweet_df['tweet'].iloc[0], "\n")
```

user father dysfunctional selfish drags kids dysfunction run

```
[60]: tweet_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 29345 entries, 0 to 31961
Data columns (total 3 columns):
#   Column  Non-Null Count  Dtype
---  -
0   id      29345 non-null    int64
1   label   29345 non-null    int64
2   tweet   29345 non-null    object
dtypes: int64(2), object(1)
memory usage: 917.0+ KB
```


```
[61]: tweet_df['label'].value_counts()
```

```
[61]: 0    27352
      1     1993
      Name: label, dtype: int64
```

3 Data Visualization

```
[62]: fig = plt.figure(figsize=(5,5))
      sns.countplot(x='label', data = tweet_df)
```


```
[62]: <Axes: xlabel='label', ylabel='count'>
```



```
[63]: fig = plt.figure(figsize=(7,7))
      colors = ("red", "gold")
      wp = {'linewidth':2, 'edgecolor':"black"}
      tags = tweet_df['label'].value_counts()
      explode = (0.1, 0.1)
      tags.plot(kind='pie',autopct = '%1.1f%%', shadow=True, colors = colors,
      ↪startangle =90,
      wedgeprops = wp, explode = explode, label='')
      plt.title('Distribution of sentiments')
```

```
[63]: Text(0.5, 1.0, 'Distribution of sentiments')
```

Distribution of sentiments




```
[64]: non_hate_tweets = tweet_df[tweet_df.label == 0]
      non_hate_tweets.head()
```

```
[64]:   id  label      tweet
0    1      0  user father dysfunctional selfish drags kids d...
1    2      0  user user thanks lyft credit cant use cause do...
2    3      0                bihday majesty
3    4      0          model love u take u time ur
4    5      0  factsguide society motivation
```

```
[65]: text = ' '.join([word for word in non_hate_tweets['tweet']])
      plt.figure(figsize=(20,15), facecolor='None')
      wordcloud = WordCloud(max_words=500, width=1600, height=800).generate(text)
```

```
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis('off')
plt.title('Frequent words in non hate tweets', fontsize = 19)
plt.show()
```

Frequent words in non hate tweets



```
[66]: hate_tweets = tweet_df[tweet_df.label == 1]
      hate_tweets.head()
```

```
[66]:      id  label      tweet
      13  14      1  user cnn calls michigan middle school build wa...
      14  15      1  comment australia opkillingbay seashepherd hel...
      17  18      1      retweet agree
      23  24      1      user user lumpy says prove lumpy
      34  35      1  unbelievable 21st century wed need something l...
```

```
[67]: text = ' '.join([word for word in hate_tweets['tweet']])
plt.figure(figsize=(20,15), facecolor='None')
wordcloud = WordCloud(max_words=500, width=1600, height=800).generate(text)
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis('off')
plt.title('Frequent words in hate tweets', fontsize = 19)
plt.show()
```

4 Text Representation(Vectorization)

- Use of Tfidf -> Text frequency - Inverse document frequency using sklearn

```
[68]: vect = TfidfVectorizer(ngram_range=(1,2)).fit(tweet_df['tweet'])
```

```
[69]: feature_names = vect.get_feature_names_out()
print("Number of features: {}".format(len(feature_names)))
print("First 10 features: {}".format(feature_names[:10]))
```

Number of features: 208579

First 10 features:

```
[ '0000001' '0000001 polluting' '00027' '00027 photooftheday' '001' '0035'
  '00h30' '01' '01 4995' '01 7900' ]
```

```
[70]: vect = TfidfVectorizer(ngram_range=(1,3)).fit(tweet_df['tweet'])
```

```
[71]: feature_names = vect.get_feature_names_out()
print("Number of features: {}".format(len(feature_names)))
print("First 10 features: \n{}".format(feature_names[:10]))
```

Number of features: 380305

First 10 features:

```
['0000001' '0000001 polluting' '0000001 polluting niger' '00027'
 '00027 photooftheday' '00027 photooftheday music' '001' '0035' '00h30'
 '01']
```

5 Model building

```
[72]: X = tweet_df['tweet']
      Y = tweet_df['label']
      X = vect.transform(X)

      x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.2,
      ↪random_state=42)
```

```
[73]: print("Size of x_train:", (x_train.shape))
      print("Size of y_train:", (y_train.shape))
      print("Size of x_test: ", (x_test.shape))
      print("Size of y_test: ", (y_test.shape))
```

Size of x_train: (23476, 380305)

Size of y_train: (23476,)

Size of x_test: (5869, 380305)

Size of y_test: (5869,)

6 Hyperparameter tuning using GridSearchCV

```
[74]: from sklearn.model_selection import GridSearchCV
      import warnings
      warnings.filterwarnings('ignore')
```

```
[75]: param_grid = {'C':[100, 10, 1.0, 0.1, 0.01], 'solver' :['newton-cg',
      ↪'lbfgs','liblinear']}
      grid_logit = GridSearchCV(LogisticRegression(), param_grid, cv = 5)
      grid_logit.fit(x_train, y_train)
      print("Best Cross validation score: {:.2f}".format(grid_logit.best_score_))
      print("Best parameters: ", grid_logit.best_params_)
```

Best Cross validation score: 0.95

Best parameters: {'C': 100, 'solver': 'newton-cg'}

```
[76]: logreg = LogisticRegression(C=100,solver= 'newton-cg')
      logreg.fit(x_train, y_train)
      logreg_predict = logreg.predict(x_test)
      logreg_acc = accuracy_score(logreg_predict, y_test)
      print("Test accuracy: {:.2f}%".format(logreg_acc*100))
```

Test accuracy: 94.89%

```
[77]: print(confusion_matrix(y_test, logreg_predict))
      print("\n")
      print(classification_report(y_test, logreg_predict))
```



```

style.use('classic')
cm = confusion_matrix(y_test, logreg_predict, labels=logreg.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=logreg.
    ↳classes_)
disp.plot()

```

```

[[5450   8]
 [ 292 119]]


```

	precision	recall	f1-score	support
0	0.95	1.00	0.97	5458
1	0.94	0.29	0.44	411
accuracy			0.95	5869
macro avg	0.94	0.64	0.71	5869
weighted avg	0.95	0.95	0.94	5869

```

[77]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
      0x7b831efa7100>

```



```
[78]: des = DecisionTreeClassifier()
des.fit(x_train, y_train)
des_predict = des.predict(x_test)
des_acc = accuracy_score(des_predict, y_test)
print("Test accuracy: {:.2f}%".format(des_acc*100))
```

Test accuracy: 94.50%

```
[79]: print(des.tree_.max_depth)
```

1190

```
[80]: des2 = DecisionTreeClassifier(criterion='entropy')
des2.fit(x_train, y_train)
des2_predict = des2.predict(x_test)
des2_acc = accuracy_score(des2_predict, y_test)
print("Test accuracy: {:.2f}%".format(des2_acc*100))
```

Test accuracy: 94.14%

```
[81]: print(des2.tree_.max_depth)
```

1102

```
[82]: svc = SVC()
      svc.fit(x_train, y_train)
      svc_predict = svc.predict(x_test)
      svc_acc = accuracy_score(svc_predict, y_test)
      print("Test accuracy: {:.2f}%".format(svc_acc*100))
```

Test accuracy: 93.99%

```
[83]: knn=KNeighborsClassifier()
      knn.fit(x_train,y_train)
      knn_predict=knn.predict(x_test)
      knn_acc = accuracy_score(knn_predict,y_test)
      print("Test accuracy: {:.2f}%".format(knn_acc*100))
```

Test accuracy: 94.43%

```
[84]: print(logreg.predict(x_test[20:30]))
      print(des.predict(x_test[20:30]))
      print(des2.predict(x_test[20:30]))
      print(svc.predict(x_test[20:30]))
      print(knn.predict(x_test[20:30]))
      print("Actual",y_test[20:30].values)
```

[0 0 0 0 1 0 0 0 0 0]

[0 0 0 0 1 0 0 0 0 0]

[0 0 0 0 0 0 0 0 0 0]

[0 0 0 0 0 0 0 0 0 0]

[0 0 0 0 0 0 0 0 0 0]

Actual [0 0 0 0 1 0 0 1 0 0]

```
[91]: def GetHateOrNonhate(text):
      text = data_processing(text)
      text = lemmatizing(text)
      text = vect.transform([text])
      pred = logreg.predict(text)
      if(pred):
          print("Hate")
      else:
          print("Non Hate")

      s = 'I love my cunning friend.'
      t = 'Africans is Black people!'
```

```
GetHateOrNonhate(s)  
GetHateOrNonhate(t)
```

Non Hate
Hate