

Hate Speech detection

Dataset: <https://www.kaggle.com/datasets/arkhoshghalb/twitter-sentiment-analysis-hatred-speech>

Importing Required Libraries

```
import pandas as pd
import numpy as np
import re
import seaborn as sns
import matplotlib.pyplot as plt
from matplotlib import style
style.use('ggplot')
from nltk.tokenize import word_tokenize
from nltk.stem import WordNetLemmatizer
from nltk.corpus import stopwords
import nltk
nltk.download('stopwords')
nltk.download('punkt')
nltk.download('wordnet')
stop_words = set(stopwords.words('english'))
from wordcloud import WordCloud
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix, ConfusionMatrixDisplay
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Package stopwords is already up-to-date!
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data]   Package punkt is already up-to-date!
[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data]   Package wordnet is already up-to-date!
```

```
tweet_df = pd.read_csv('/content/train.csv')
tweet_df.head()
```

	id	label	tweet
0	1	0	@user when a father is dysfunctional and is s...
1	2	0	@user @user thanks for #lyft credit i can't us...
2	3	0	bihday your majesty
3	4	0	#model i love u take with u all the time in ...
4	5	0	factsguide: society now #motivation

```

tweet_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 31962 entries, 0 to 31961
Data columns (total 3 columns):
#   Column  Non-Null Count  Dtype
---  -
0    id      31962 non-null    int64
1   label    31962 non-null    int64
2   tweet    31962 non-null    object
dtypes: int64(2), object(1)
memory usage: 749.2+ KB

print(tweet_df['tweet'].iloc[0], "\n")

@user when a father is dysfunctional and is so selfish he drags his
kids into his dysfunction.    #run

```

Data Preprocessing

- Removing links, mails, more than one spaces
- Tokenization using NLTK
- Removing stop words

```

def data_processing(tweet):
    tweet = tweet.lower()
    tweet = re.sub(r"https\S+|www\S+http\S+", '', tweet, flags =
re.MULTILINE)
    tweet = re.sub(r'\@w+|\#','', tweet)
    tweet = re.sub(r'^\w\s','', tweet)
    tweet = re.sub(r'ö','', tweet)
    tweet_tokens = word_tokenize(tweet)
    filtered_tweets = [w for w in tweet_tokens if not w in stop_words]
    return " ".join(filtered_tweets)

tweet_df.tweet = tweet_df['tweet'].apply(data_processing)
tweet_df = tweet_df.drop_duplicates('tweet')

```

- Lemmatization Using NLTK

```

lemmatizer = WordNetLemmatizer()
def lemmatizing(data):
    tweet = [lemmatizer.lemmatize(word) for word in data]
    return data

tweet_df['tweet'] = tweet_df['tweet'].apply(lambda x: lemmatizing(x))
print(tweet_df['tweet'].iloc[0], "\n")

```

```
user father dysfunctional selfish drags kids dysfunction run
```

```
tweet_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 29345 entries, 0 to 31961
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  -
0   id           29345 non-null  int64
1   label        29345 non-null  int64
2   tweet        29345 non-null  object
dtypes: int64(2), object(1)
memory usage: 917.0+ KB
```

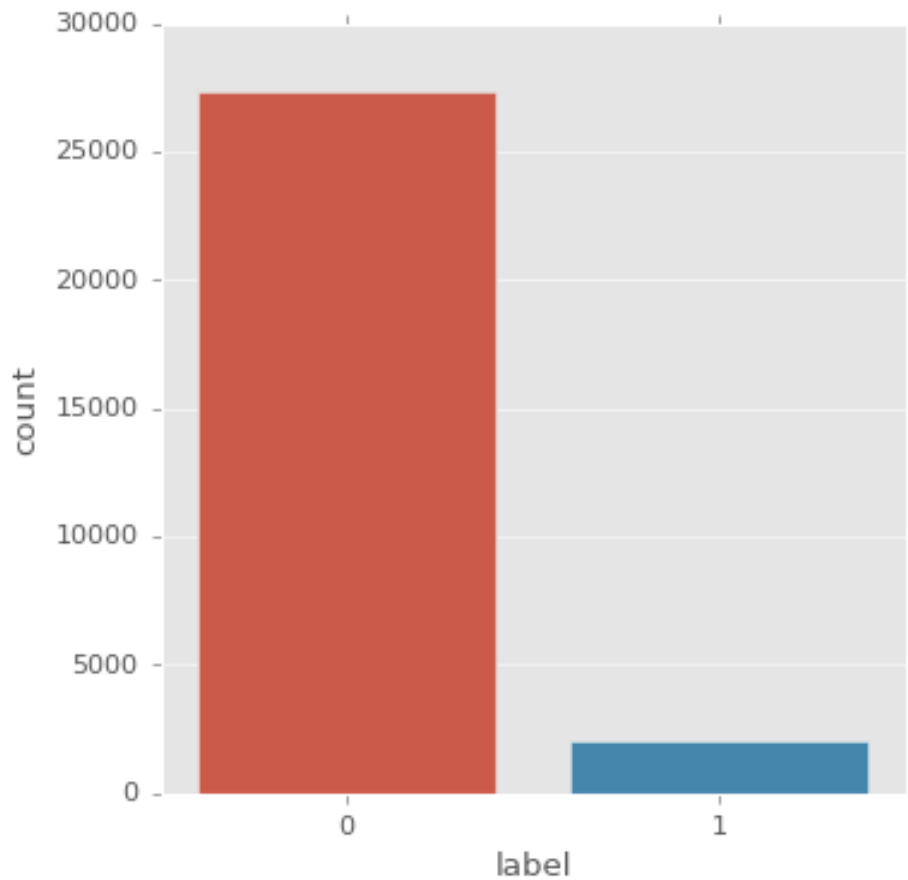
```
tweet_df['label'].value_counts()
```

```
0    27352
1     1993
Name: label, dtype: int64
```

Data Visualization


```
fig = plt.figure(figsize=(5,5))
sns.countplot(x='label', data = tweet_df)
```

```
<Axes: xlabel='label', ylabel='count'>
```



```
fig = plt.figure(figsize=(7,7))
colors = ("red", "gold")
wp = {'linewidth':2, 'edgecolor':"black"}
tags = tweet_df['label'].value_counts()
explode = (0.1, 0.1)
tags.plot(kind='pie',autopct = '%1.1f%%', shadow=True, colors =
colors, startangle =90,
        wedgeprops = wp, explode = explode, label='')
plt.title('Distribution of sentiments')
Text(0.5, 1.0, 'Distribution of sentiments')
```

Distribution of sentiments




```
non_hate_tweets = tweet_df[tweet_df.label == 0]
non_hate_tweets.head()
```

	id	label	tweet
0	1	0	user father dysfunctional selfish drags kids d...
1	2	0	user user thanks lyft credit cant use cause do...
2	3	0	bihday majesty
3	4	0	model love u take u time ur
4	5	0	factsguide society motivation

```
text = ' '.join([word for word in non_hate_tweets['tweet']])
plt.figure(figsize=(20,15), facecolor='None')
wordcloud = WordCloud(max_words=500, width=1600,
height=800).generate(text)
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis('off')
```

```
plt.title('Frequent words in non hate tweets', fontsize = 19)
plt.show()
```

Frequent words in non hate tweets



```
hate_tweets = tweet_df[tweet_df.label == 1]
hate_tweets.head()
```

	id	label	tweet
13	14	1	user cnn calls michigan middle school build wa...
14	15	1	comment australia opkillingbay seashepherd hel...
17	18	1	retweet agree
23	24	1	user user lumpy says prove lumpy
34	35	1	unbelievable 21st century wed need something l...

```
text = ' '.join([word for word in hate_tweets['tweet']])
plt.figure(figsize=(20,15), facecolor='None')
wordcloud = WordCloud(max_words=500, width=1600,
height=800).generate(text)
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis('off')
plt.title('Frequent words in hate tweets', fontsize = 19)
plt.show()
```


- Use of Tfidf -> Text frequency - Inverse document frequency using sklearn

- ```
vect = TfidfVectorizer(ngram_range=(1,2)).fit(tweet_df['tweet'])
```

Number of features: 208579

```
vect = TfidfVectorizer(ngram_range=(1,3)).fit(tweet_df['tweet'])
```

Number of features: 380305

```
First 10 features:
['0000001' '0000001 polluting' '0000001 polluting niger' '00027'
 '00027 photooftheday' '00027 photooftheday music' '001' '0035'
 '00h30'
 '01']
```

# Model building

```
X = tweet_df['tweet']
Y = tweet_df['label']
X = vect.transform(X)

x_train, x_test, y_train, y_test = train_test_split(X, Y,
test_size=0.2, random_state=42)

print("Size of x_train:", (x_train.shape))
print("Size of y_train:", (y_train.shape))
print("Size of x_test: ", (x_test.shape))
print("Size of y_test: ", (y_test.shape))

Size of x_train: (23476, 380305)
Size of y_train: (23476,)
Size of x_test: (5869, 380305)
Size of y_test: (5869,)
```

# Hyperparameter tuning using GridSearchCV

```
from sklearn.model_selection import GridSearchCV
import warnings
warnings.filterwarnings('ignore')

param_grid = {'C':[100, 10, 1.0, 0.1, 0.01], 'solver':['newton-cg',
'lbfgs','liblinear']}
grid_logit = GridSearchCV(LogisticRegression(), param_grid, cv = 5)
grid_logit.fit(x_train, y_train)
print("Best Cross validation score:
{:.2f}".format(grid_logit.best_score_))
print("Best parameters: ", grid_logit.best_params_)

Best Cross validation score: 0.95
Best parameters: {'C': 100, 'solver': 'newton-cg'}

logreg = LogisticRegression(C=100,solver= 'newton-cg')
logreg.fit(x_train, y_train)
logreg_predict = logreg.predict(x_test)
logreg_acc = accuracy_score(logreg_predict, y_test)
print("Test accuracy: {:.2f}%".format(logreg_acc*100))

Test accuracy: 94.89%

print(confusion_matrix(y_test, logreg_predict))
print("\n")
print(classification_report(y_test, logreg_predict))

style.use('classic')
```




```
cm = confusion_matrix(y_test, logreg_predict, labels=logreg.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix=cm,
display_labels=logreg.classes_)
disp.plot()
```

```
[[5450 8]
 [292 119]]
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.95      | 1.00   | 0.97     | 5458    |
| 1            | 0.94      | 0.29   | 0.44     | 411     |
| accuracy     |           |        |          | 5869    |
| macro avg    |           |        |          | 5869    |
| weighted avg |           |        |          | 5869    |

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x7ade1f4de290>
```



```
des = DecisionTreeClassifier()
des.fit(x_train, y_train)
des_predict = des.predict(x_test)
des_acc = accuracy_score(des_predict, y_test)
print("Test accuracy: {:.2f}%".format(des_acc*100))

Test accuracy: 94.45%

des2 = DecisionTreeClassifier(criterion='entropy')
des2.fit(x_train, y_train)
des2_predict = des2.predict(x_test)
des2_acc = accuracy_score(des2_predict, y_test)
print("Test accuracy: {:.2f}%".format(des2_acc*100))

Test accuracy: 93.83%

svc = SVC()
svc.fit(x_train, y_train)
svc_predict = svc.predict(x_test)
```

```
svc_acc = accuracy_score(svc_predict, y_test)
print("Test accuracy: {:.2f}%".format(svc_acc*100))
```

Test accuracy: 93.99%

```
knn=KNeighborsClassifier()
knn.fit(x_train,y_train)
knn_predict=knn.predict(x_test)
knn_acc = accuracy_score(knn_predict,y_test)
print("Test accuracy: {:.2f}%".format(knn_acc*100))
```

Test accuracy: 94.43%

```
print(logreg.predict(x_test[20:30]))
print(des.predict(x_test[20:30]))
print(des2.predict(x_test[20:30]))
print(svc.predict(x_test[20:30]))
print(knn.predict(x_test[20:30]))
print("Actual",y_test[20:30].values)
```

```
[0 0 0 0 1 0 0 0 0 0]
[0 0 0 0 1 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0]
```

Actual [0 0 0 0 1 0 0 1 0 0]